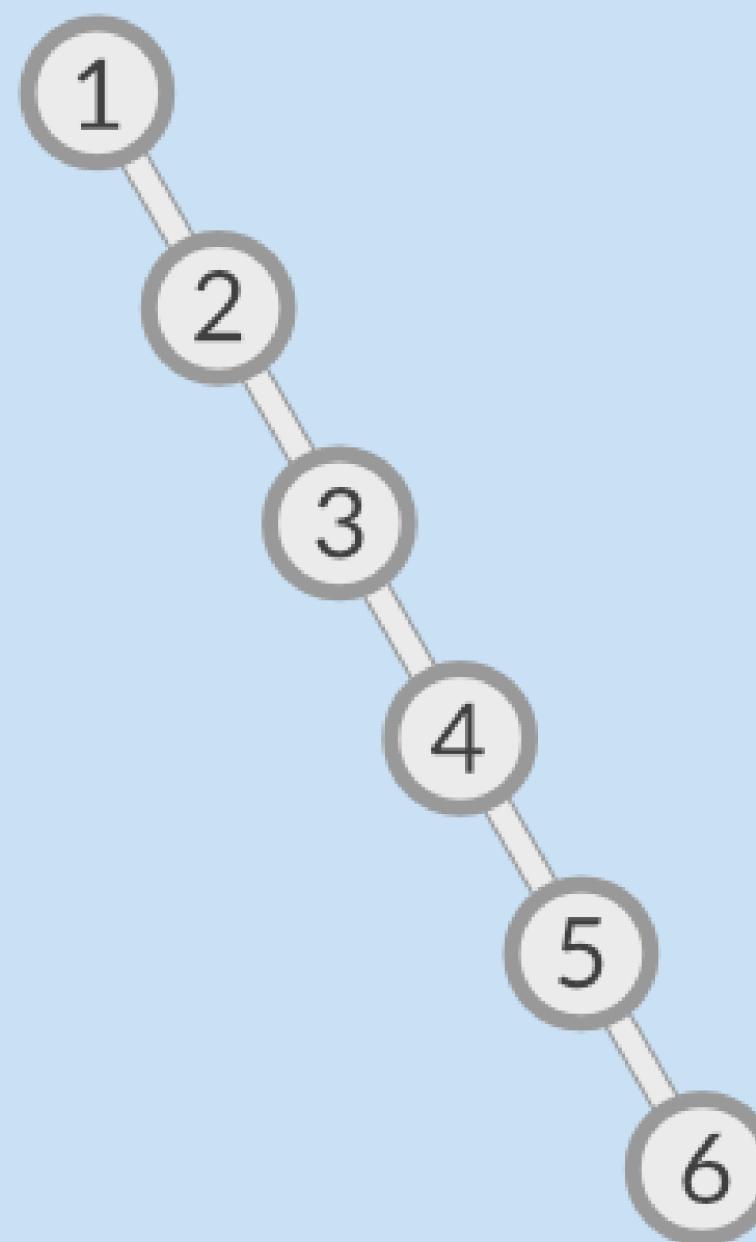
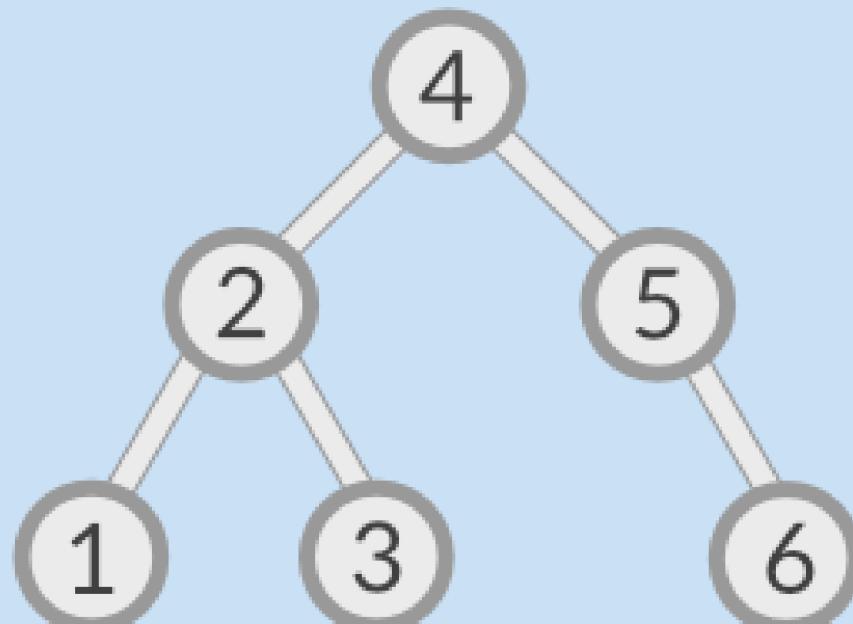


UNDERSTANDING RED-BLACK TREES



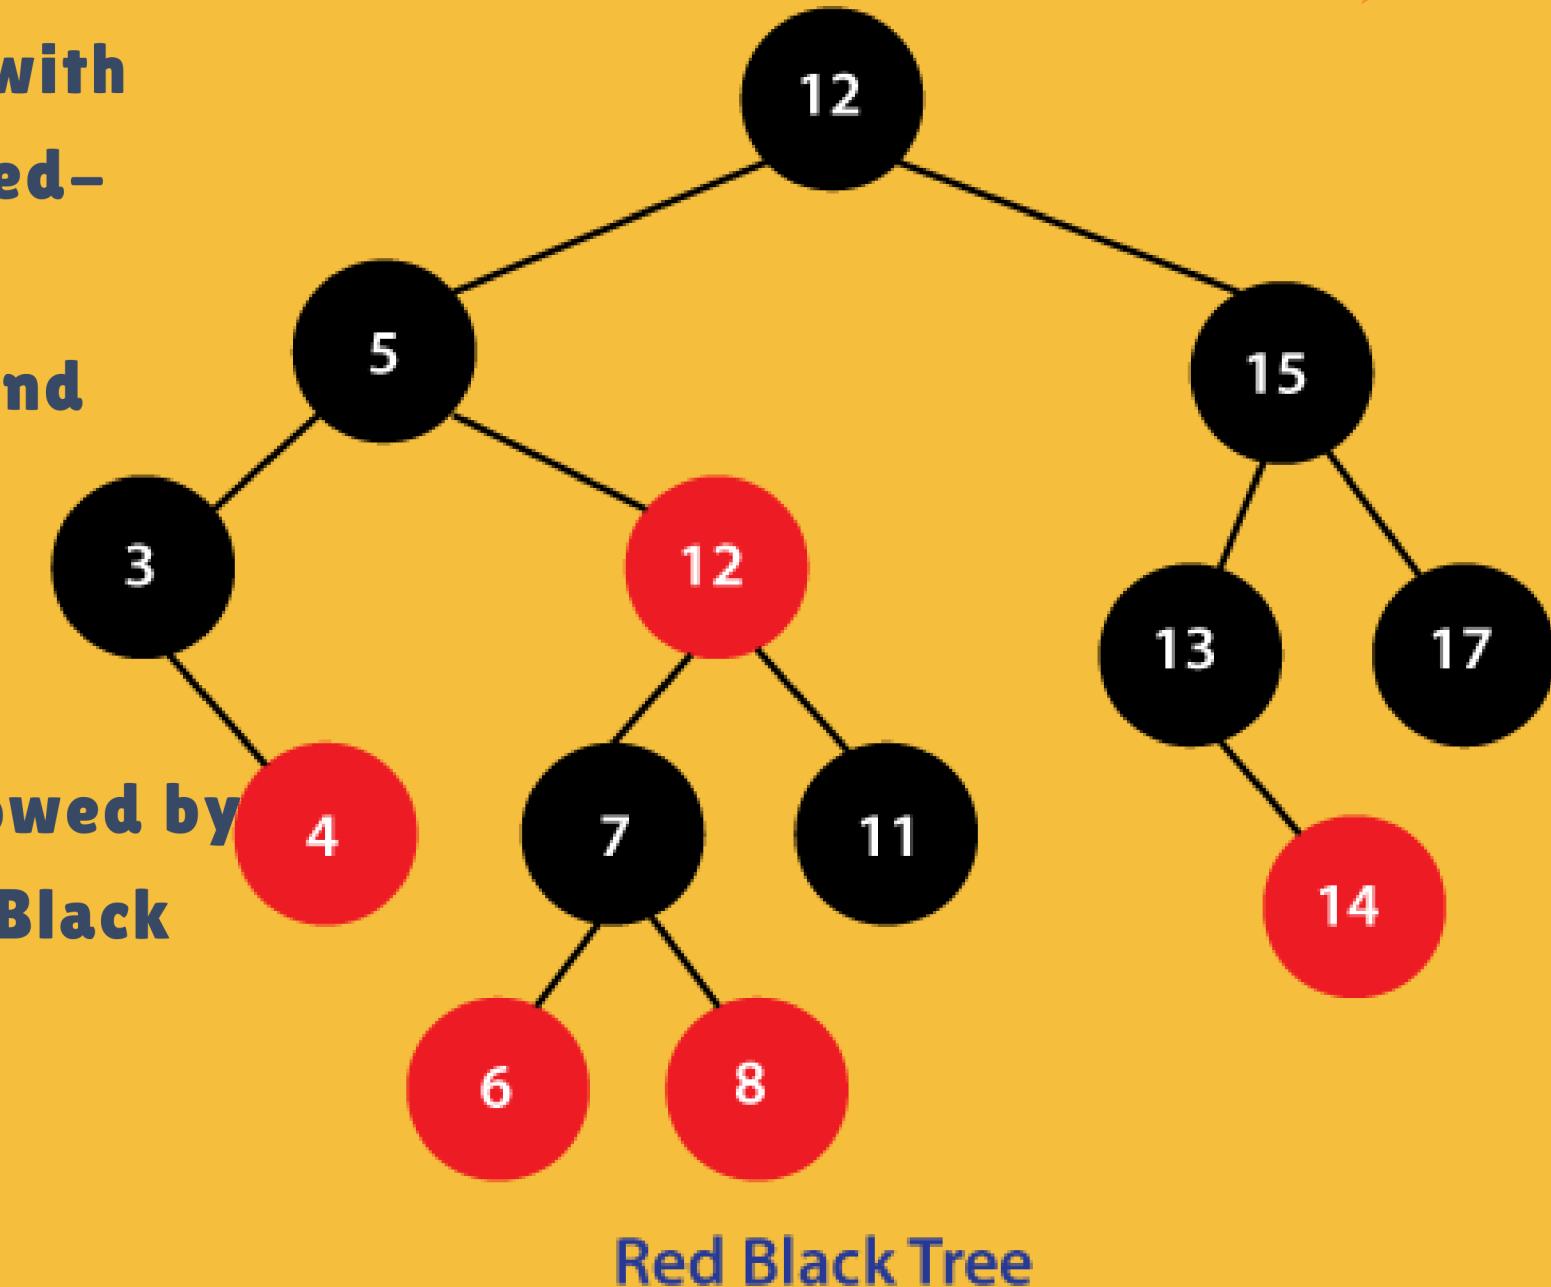
THE NEED FOR BALANCE IN TREES



- A balanced BST offers the best search efficiency because the path from the root to any leaf node is of minimal length, which ideally is $O(\log n)$.
- An unbalanced tree can degenerate into a linked list structure in the worst case, making the search time complexity linear, or $O(n)$.

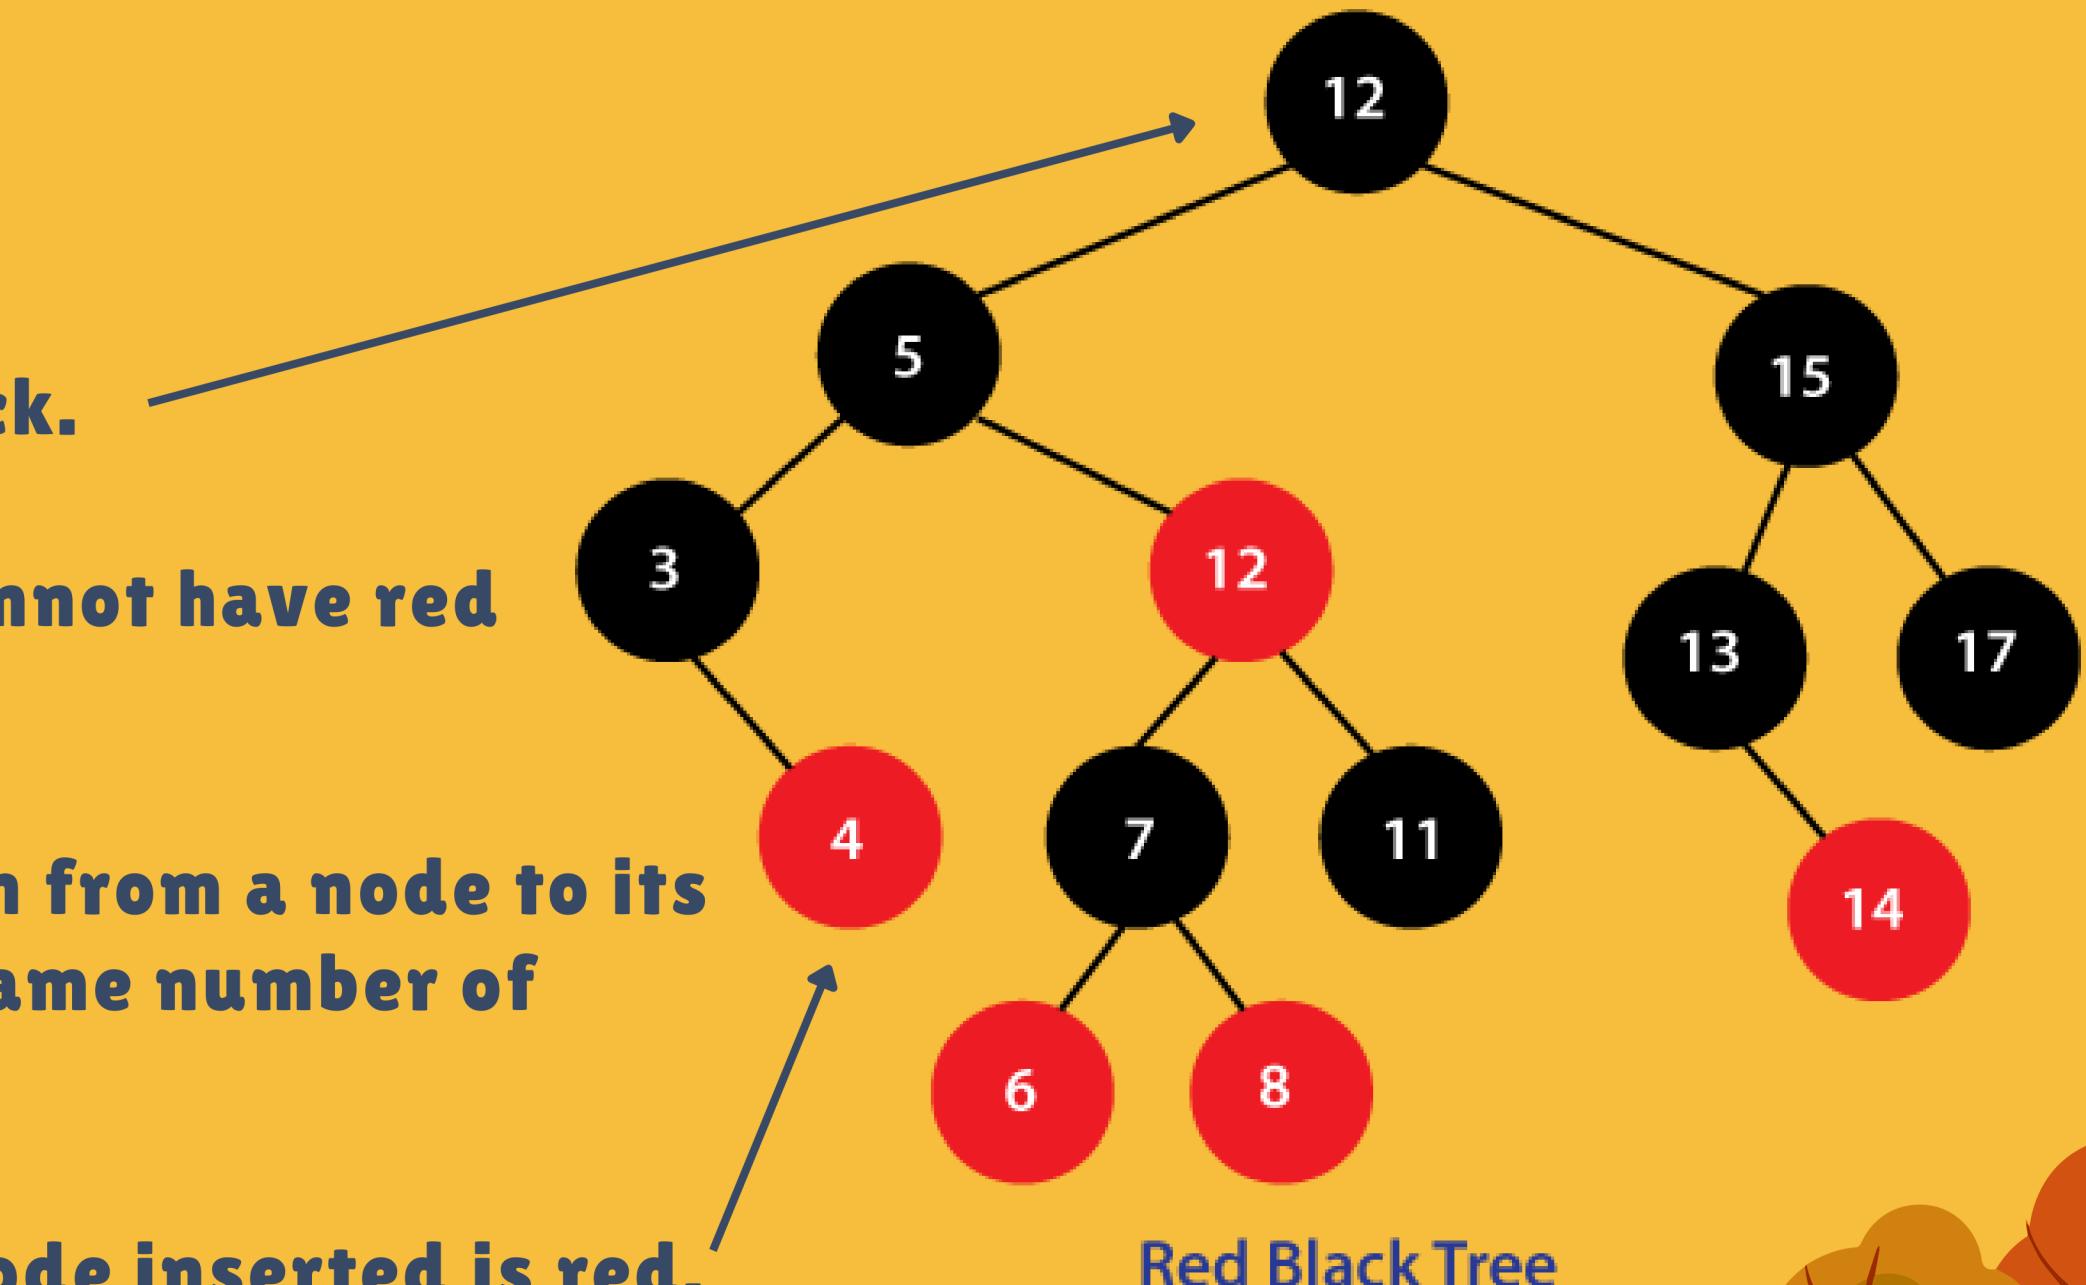
RED-BLACK TREE - AN OVERVIEW

- A Red-Black Tree is a type of binary search tree with extra properties that lead to a balanced tree. A Red-Black Tree is self-balancing, which means it automatically rearranges itself after insertions and deletions to maintain balance.
- Rotation Rules: Insertions and deletions are followed by rotations and color changes to maintain the Red-Black properties.



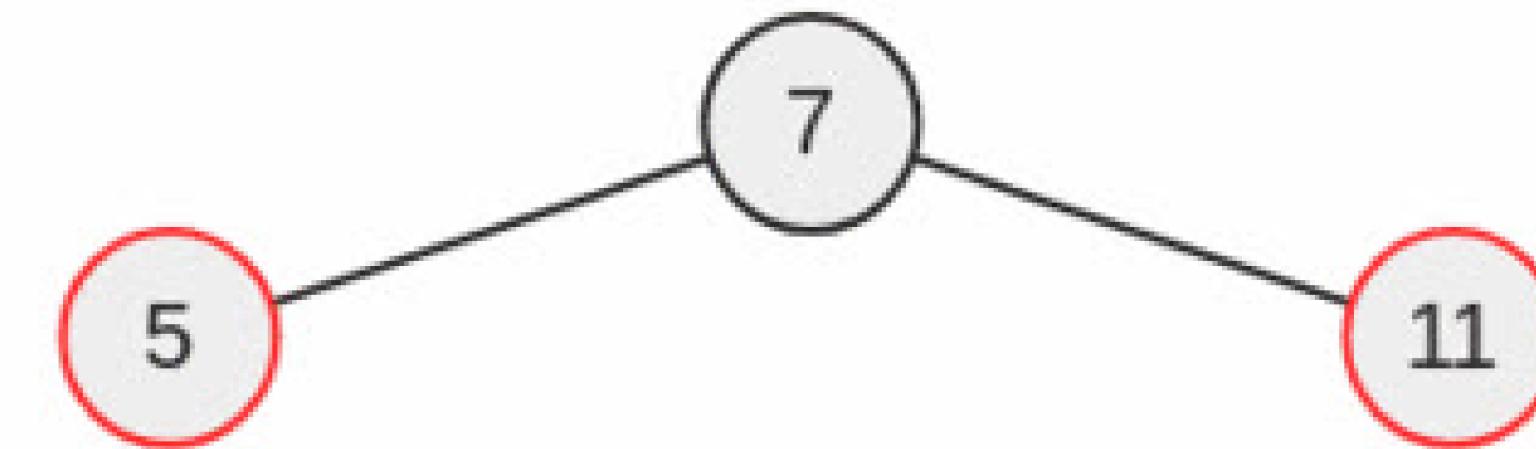
RED-BLACK TREE PROPERTIES

- Every node is either red or black.
- All leaves (NIL nodes) are black.
- The root of the tree is always black.
- Red node property: Red nodes cannot have red children.
- Black Height Property: Every path from a node to its descendant NULL nodes has the same number of black nodes.
- New Node Property: Every new node inserted is red.



INSERTION IN A RED-BLACK TREE

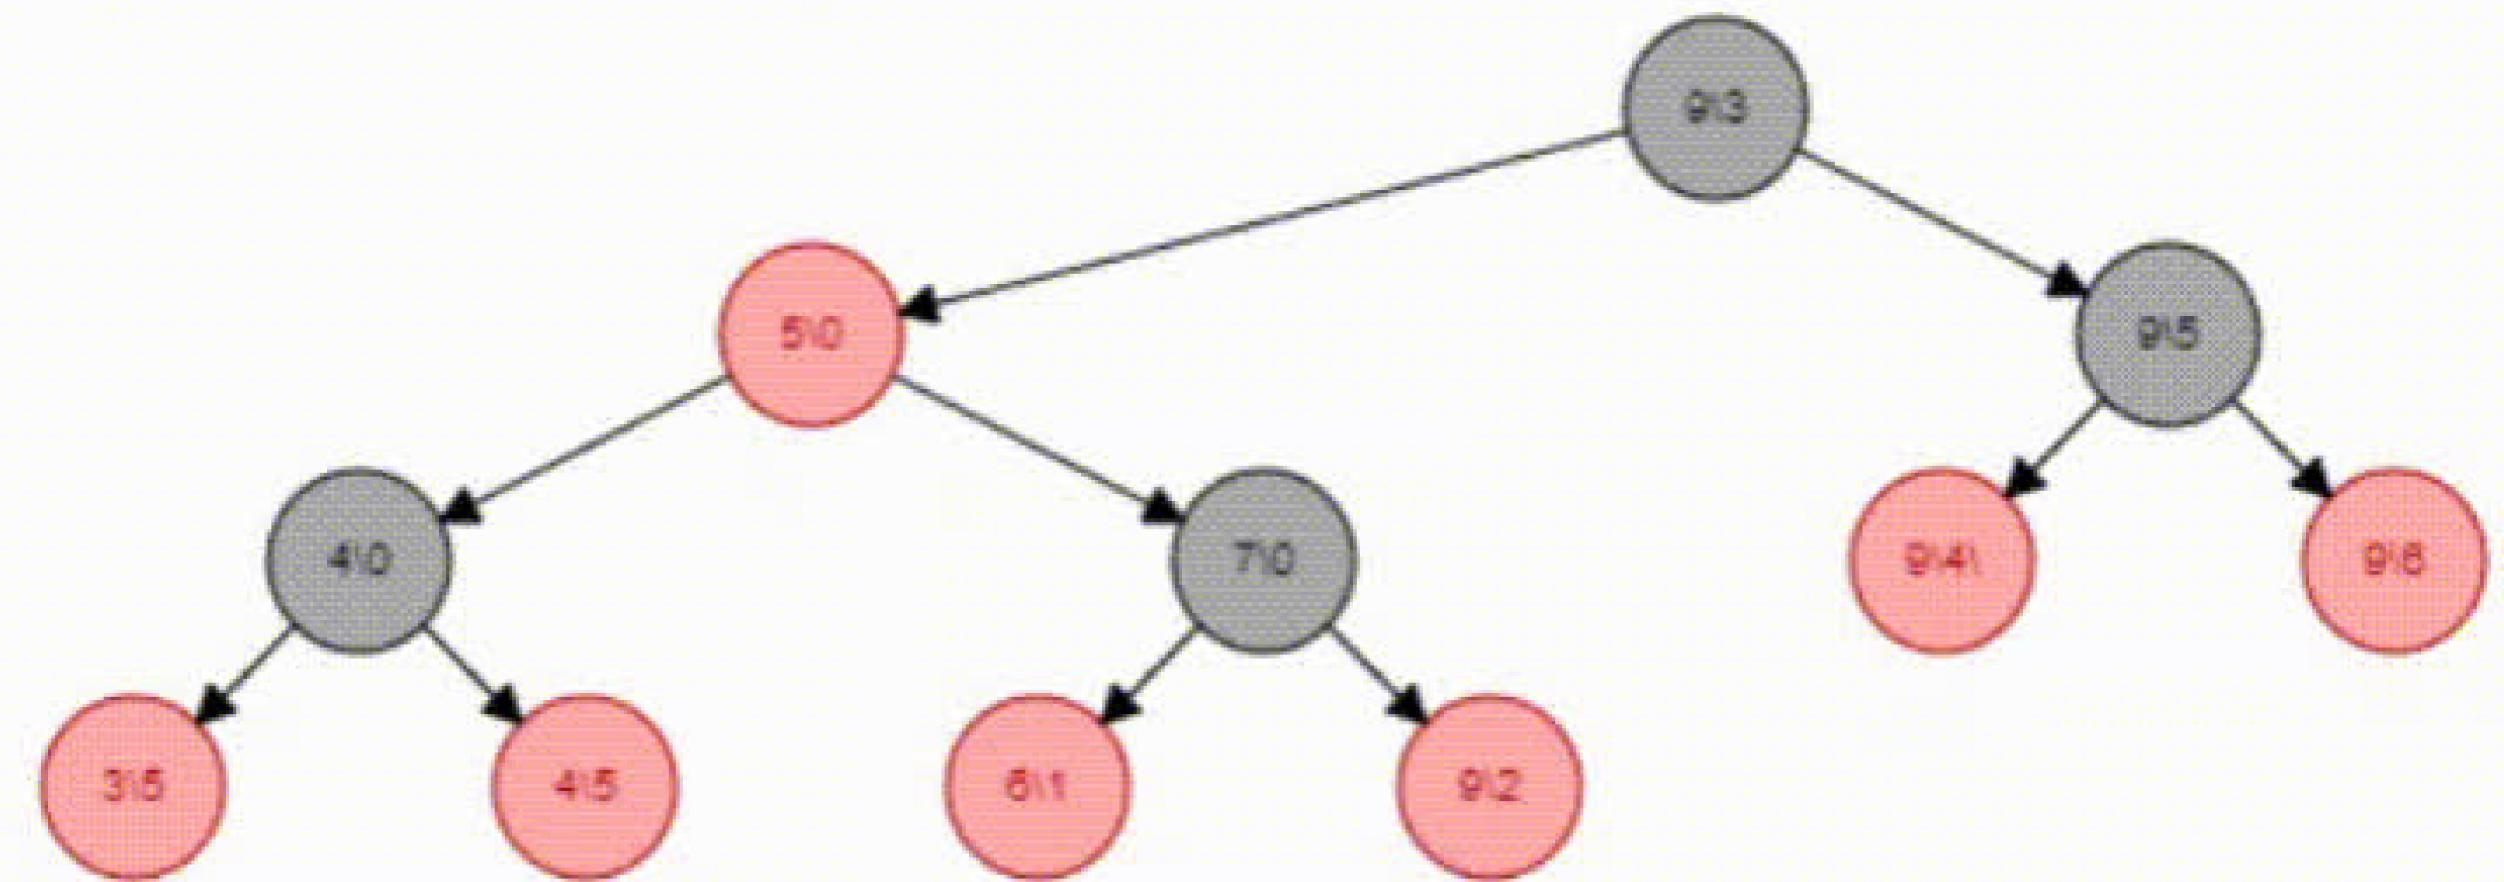
- **Similarity to BST:** Insertion begins just like in a regular Binary Search Tree, by placing the new node in the correct position according to its value.
- **Initial Coloring:** The new node is always colored red at first, to maintain the black-height property.



DELETION IN A RED-BLACK TREE

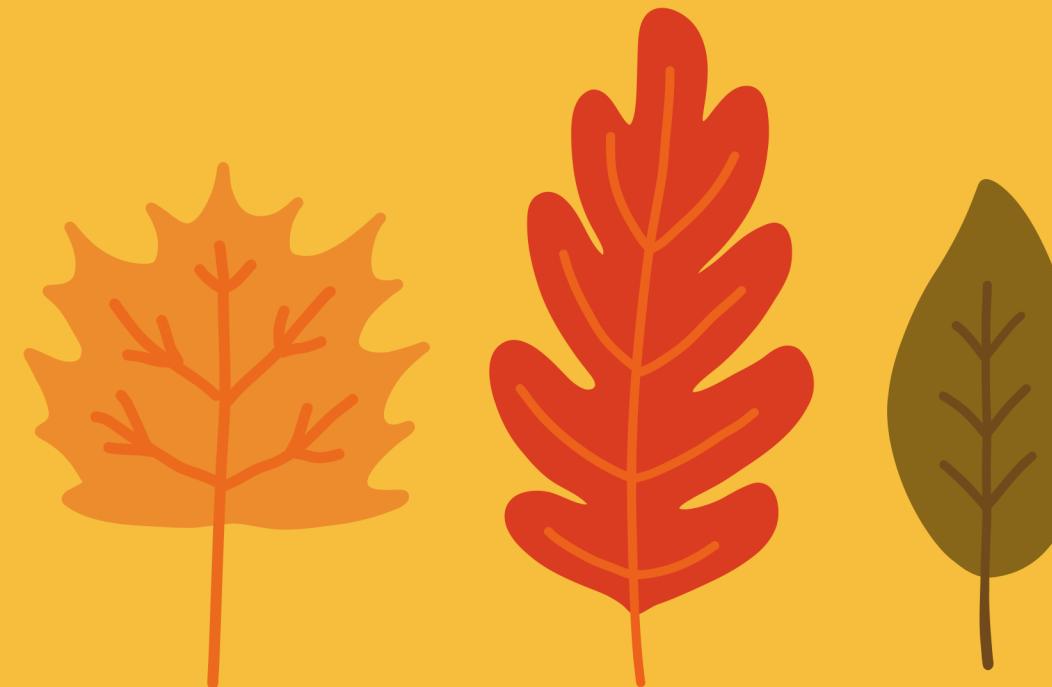
- Preliminary Note: Deletion in a Red–Black Tree can be complex because it may disrupt the tree's balancing properties.

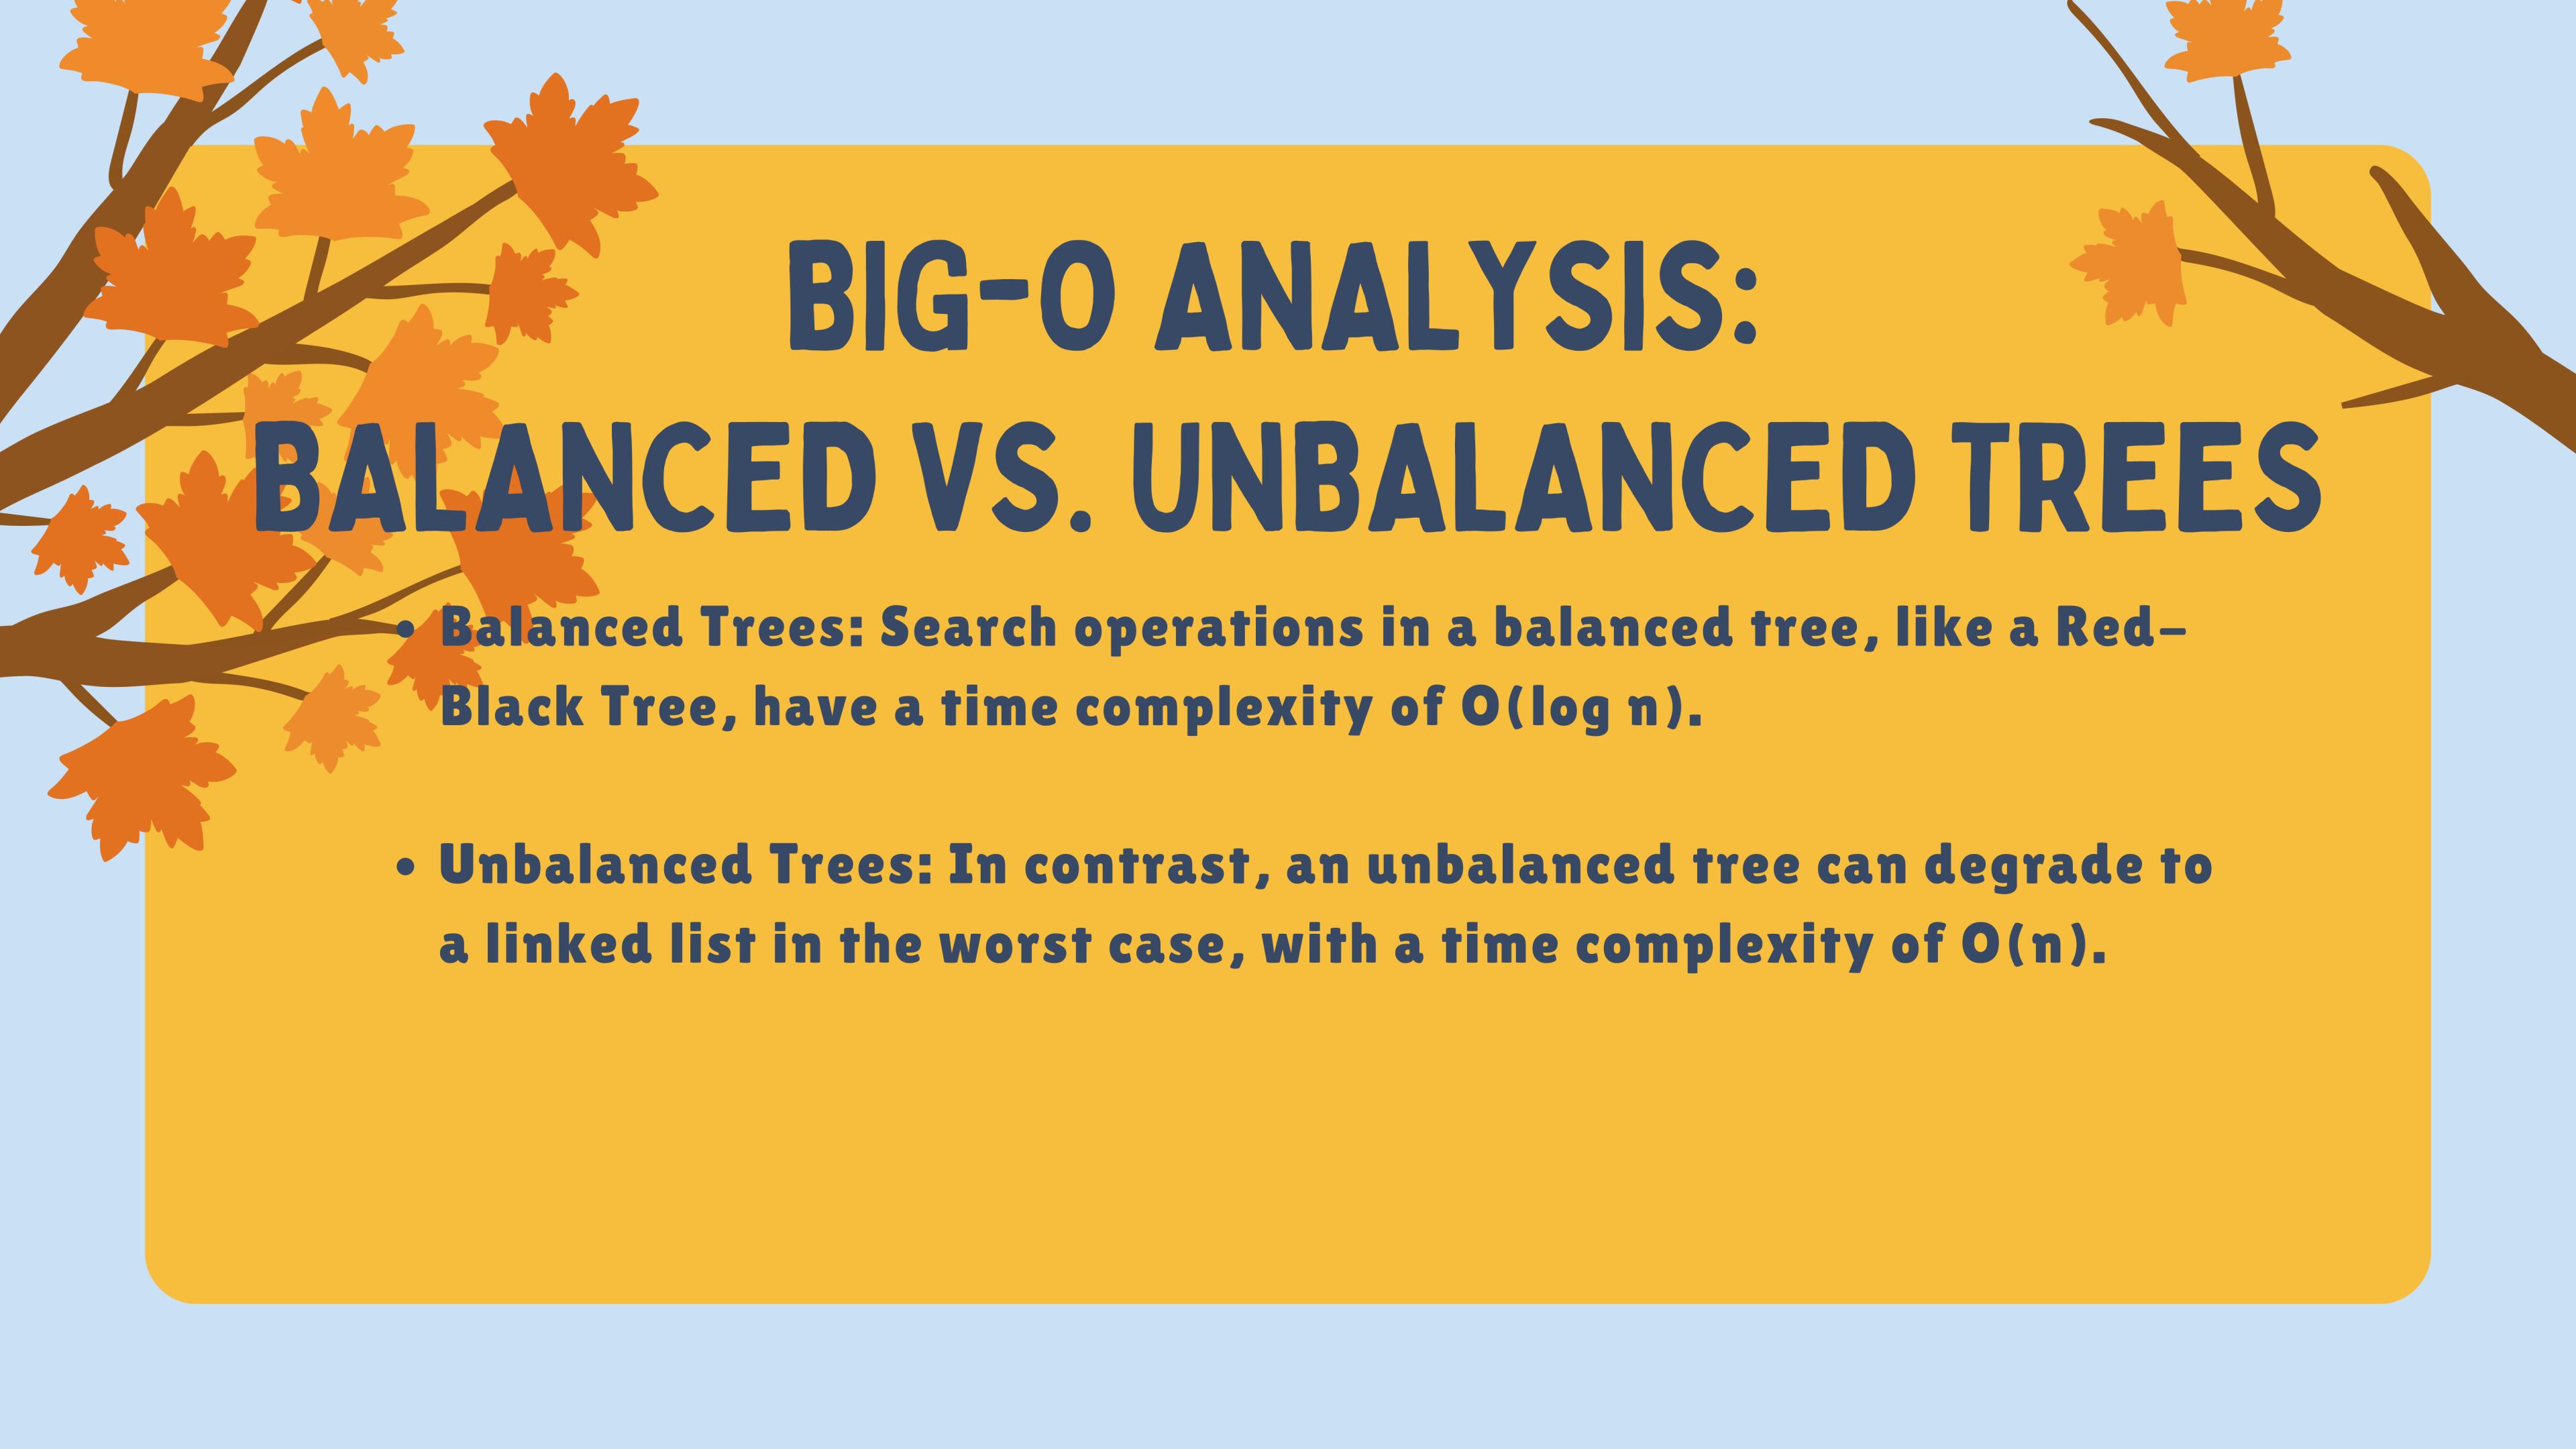
- Initial Step: Just like in a Binary Search Tree, we locate the node to be deleted first.



ADVANTAGES OF BALANCED TREES

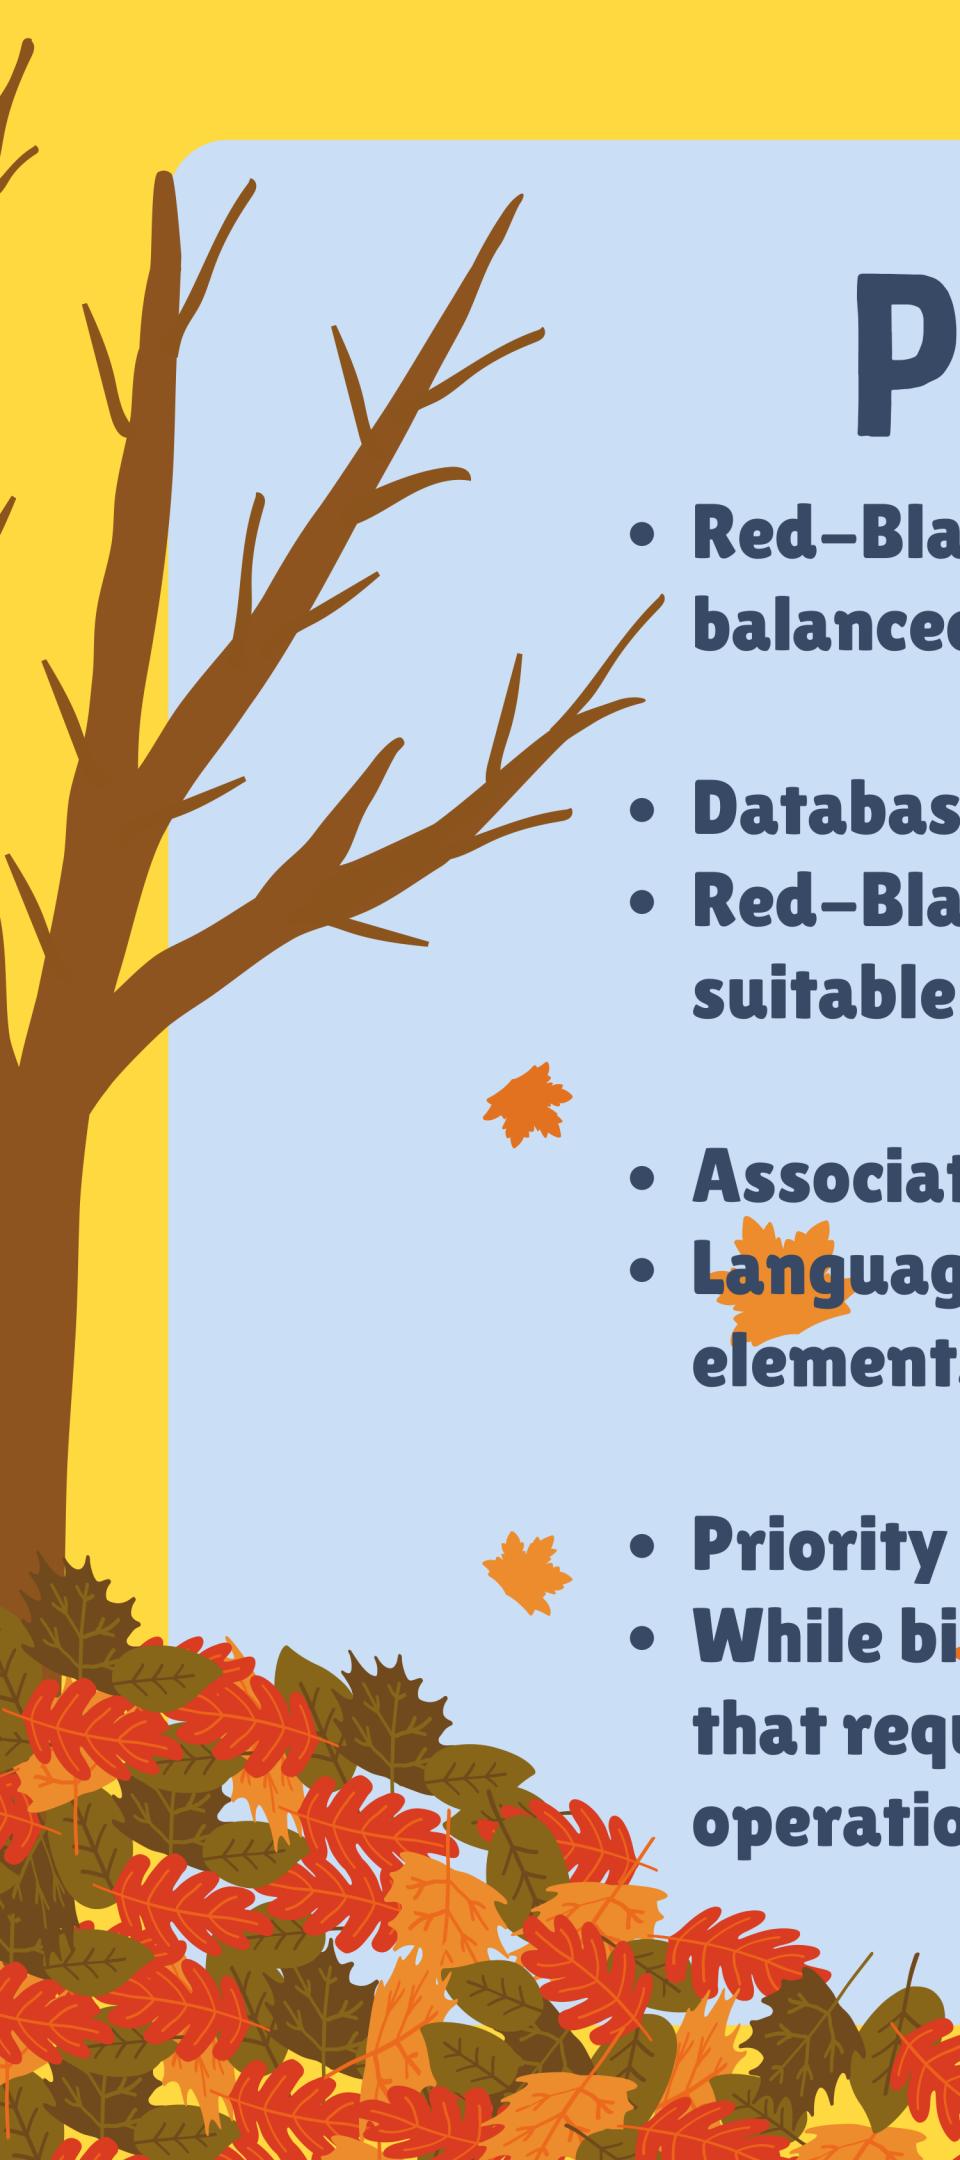
- Efficient search, insertion, and deletion operations
- Improved performance in dynamic, real-world data usage scenarios
- Ensured $\log(n)$ height for balanced data distribution





BIG-O ANALYSIS: BALANCED VS. UNBALANCED TREES

- **Balanced Trees:** Search operations in a balanced tree, like a Red-Black Tree, have a time complexity of $O(\log n)$.
- **Unbalanced Trees:** In contrast, an unbalanced tree can degrade to a linked list in the worst case, with a time complexity of $O(n)$.



PRACTICAL APPLICATION

- Red–Black Trees are particularly effective in scenarios where we need to maintain a balanced tree structure after frequent insertions and deletions.
- Databases:
- Red–Black Trees provide quick search, insert, and delete operations, making them suitable for database indexes where balance is crucial for performance.
- Associative Arrays:
- Languages like C++ use Red–Black Trees in implementing ordered maps, which require elements to be sorted and allow for efficient retrieval by key.
- Priority Queues:
- While binary heaps are more common, Red–Black Trees can be used for priority queues that require an ordered sequence of elements and support for other associative operations.



CONCLUSION

Red-Black Trees are a testament to the power of balanced tree structures, offering adaptability and efficiency, which makes them suitable for a wide array of applications where balanced search times are crucial.

